

Deep Reinforcement Learning Nanodegree Links
Course Description
Knowledge
Forums
DRLND Office Hours Calendar
Textbook: Reinforcement Learning: An Introduction - second edition - by Richard S. Sutton and Andrew G. Barto
github
Special Topics: Dynamic Programming
DRLND Leaderboard
openai / gym Leaderboard
Waffle (issues)
ZenDesk
Site Status updates
Deadlines (P1: Aug 28, P2: Oct 16, P3: Oct 30, Term ends: Oct 30)
Project 1 rubric
Project 2 rubric
Sections below: Articles Blogs Books Cartoons Cheatsheets Cloud Conferences Community Competitions Courses
Articles
Reinforcement Learning Doesn't Work Yet.
Why RL is flawed
How to fix RL
Evolution Strategies as a Scalable Alternative to Reinforcement Learning
Evolutionary algorithm outperforms deep-learning machines at video games
Reinforcement Learning or Evolutionary Strategies? Nature has a solution: Both.
Metacar
The Essential Guide to Training Data
Machine Learning for Humans
Math for Deep Learning
Dreaming about Driving
How To Learn Data Science If You're Broke
Beyond DQN/A3C: A Survey in Advanced Reinforcement Learning
How to rapidly test dozens of deep learning models in Python
How the Lottery Ticket Hypothesis is Challenging Everything we Knew About Training Neural Networks
Blogs
Resources for Deep Reinforcement Learning
My Curated List of AI and Machine Learning Resources from Around the Web
DeepMind
- Open sourcing TRFL: a library of reinforcement learning building blocks
OpenAI
- Reinforcement Learning with Prediction-Based Rewards
- Spinning Up in Deep RL
Tensorflow
The Gradient
UC Berkeley AI Research
Andrej Karpathy blog (older)
Deep Reinforcement Learning: Pong from Pixels
Andrej Karpathy blog (newer)
Richard S. Sutton
Moritz Hardt
Adrian Colyer: the morning paper
Towards Data Science
- What's New in Deep Learning Research: Stronger Learning with Differentiable Plasticity
Algorithmia
Locally Optimistic: The Blacker the Box
Devan Stormont
Google AI Blog
Google Developers
- Rules of Machine Learning: Best Practices for ML Engineering
- Google Developers Launchpad introduces The Lever, sharing applied-Machine Learning best practices
The Lever
The 5 Best Machine Learning GitHub Repositories & Reddit Threads from August 2018
Synced: AI Industry & Technology Review
Facebook accelerates AI development with new partners and production capabilities for PyTorch 1.0
RL—The Math behind TRPO & PPO
Adventures in Unity ML-Agents
Books
Deep Reinforcement Learning Hands-On
Grokking Deep Reinforcement Learning
The Definitive C++ Book Guide and List
Multi-Agent Machine Learning: A Reinforcement Approach

http://bit.ly/drlndlinks	Please Insert > Row and add new links freely
https://www.udacity.com/course/deep-reinforcement-learning-nanodegree--nd893	
https://knowledge.udacity.com/	
https://discussions.udacity.com/	
https://calendar.google.com/calendar/embed?src=knowlabs.com_gu20ftpeljmm1u9evmtn8vum1k@group.calendar.google.com	
https://s3-us-west-1.amazonaws.com/udacity-drlnd/bookdraft2018.pdf	https://github.com/ShangtongZhang/reinforcement-learning
https://github.com/udacity/deep-reinforcement-learning	
https://classroom.udacity.com/nanodegrees/nd893/parts/23d1307b-b908-436f-bdfe-78b6c5712b04/modules/9765795d-fea0-40-80-80	8 Hours of Extracurricular Content here!
https://docs.google.com/spreadsheets/d/1hl3RCLm21JXmcolvomYKgyXJNT9vv2bn_iFGJQw3n7I/edit#gid=0	<<< Add your best project scores and write ups here
https://github.com/openai/gym/wiki/Leaderboard	
https://waffle.io/udacity/drlnd-issues-tracker	
https://udacity.zendesk.com/hc/en-us/requests/new	
https://twitter.com/udacity	
https://classroom.udacity.com/nanodegrees/nd893/parts/8f607726-757e-4ef5-8b64-f2368755b89a/modules/a85374fa-6a60-42-80-80	
https://review.udacity.com/#!/rubrics/1889/view	
https://review.udacity.com/#!/rubrics/1890/view	https://youtu.be/2N9EoF6pQyE
Github Glossary Infographics Jobs Notebooks Papers Patents Slides Tools Videos	
https://www.alexirpan.com/2018/02/14/rl-hard.html	
https://thegradient.pub/why-rl-is-flawed/	RL works when problem is deterministic, discrete, static,
https://thegradient.pub/how-to-fix-rl/	Combine reinforcement learning and meta learning - me
https://blog.acolyer.org/	
https://www.technologyreview.com/s/611568/evolutionary-algorithm-outperforms-deep-learning-machines-at-video-games/	
https://medium.com/beyond-intelligence/reinforcement-learning-or-evolutionary-strategies-nature-has-a-solution-both-8bc80db	
https://www.metacar-project.com/	
https://visit.figure-eight.com/rs/416-ZBE-142/images/The%20Essential%20Guide%20to%20Training%20Data.pdf	
https://www.dropbox.com/s/e38nii1dn17481q/machine_learning.pdf?dl=0	
http://leiluoray.com/2018/08/29/Deep-Learning-Math/	
https://wayve.ai/blog/dreaming-about-driving-imagination-rl	
https://towardsdatascience.com/how-to-learn-data-science-if-youre-broke-7ecc408b53c7	
https://towardsdatascience.com/advanced-reinforcement-learning-6d769f529eb3	Deep RL models are really hard to train, period.
https://towardsdatascience.com/how-to-rapidly-test-dozens-of-deep-learning-models-in-python-cb839b518531	
https://towardsdatascience.com/how-the-lottery-ticket-hypothesis-is-challenging-everything-we-knew-about-training-neural-net	
https://medium.com/@yuxili/resources-for-deep-reinforcement-learning-a5fdf2dc730f	Over 120 links ...
https://medium.com/machine-learning-in-practice/my-curated-list-of-ai-and-machine-learning-resources-from-around-the-web-9	
https://deeppmind.com/blog/	
https://deeppmind.com/blog/trfl/	https://github.com/deeppmind/trfl/
https://blog.openai.com/	see OpenAI tab
https://blog.openai.com/reinforcement-learning-with-prediction-based-rewards/	
https://blog.openai.com/spinning-up-in-deep-rl/	Includes examples of RL code, educational exercises, d
https://medium.com/@tensorflow	
https://thegradient.pub/	Stanford Artificial Intelligence Laboratory (SAIL)
http://bair.berkeley.edu/blog/	
http://karpathy.github.io/	
http://karpathy.github.io/2016/05/31/rl/	
https://medium.com/@karpathy	30K followers
http://incompleteideas.net/	
http://blog.mrtz.org/	
https://blog.acolyer.org/	
https://towardsdatascience.com/	Now featuring one of our own: Partha Pratim Neog
https://towardsdatascience.com/whats-new-in-deep-learning-research-stronger-learning-with-differentiable-plasticity-9b793a5e	
https://blog.algorithmia.com/introduction-to-reinforcement-learning/	
https://www.locallyoptimistic.com/post/the-blacker-the-box/	
https://voyageintech.com/	DRLND Student
https://ai.googleblog.com/	
https://developers.googleblog.com/	
https://developers.google.com/machine-learning/guides/rules-of-ml	
https://developers.googleblog.com/2018/08/google-developers-launchpad-introduces.html	
https://medium.com/the-lever	
https://www.analyticsvidhya.com/blog/2018/09/best-machine-learning-github-repositories-reddit-threads-august-2018/	
https://syncedreview.com/	
https://code.fb.com/ai-research/facebook-accelerates-ai-development-with-new-partners-and-production-capabilities-for-pytorc	PyTorch 1.0
https://medium.com/@jonathan_hui/rl-the-math-behind-trpo-ppo-d12f6c745f33	
http://adventuresinunitymlagents.com/	By Mike Richardson and Patrick Nalepka of DRLND
https://www.amazon.com/gp/product/1788834240/ref=oh_aui_detailpage_o00_s00?ie=UTF8&psc=1#customerReviews	
https://www.manning.com/books/grokking-deep-reinforcement-learning	https://github.com/mimoralea/gdrl
https://stackoverflow.com/questions/388242/the-definitive-c-book-guide-and-list	
https://www.amazon.com/Multi-Agent-Machine-Learning-Reinforcement-Approach/dp/111836208X	

<https://drive.google.com/file/d/1xeUDVGWGUUv1-ccUMAZHJLej2C7aAFWY/view>

<https://github.com/Unity-Technologies/ml-agents/blob/master/docs/Learning-Enviro>

fully observable, fully-known, single-agent, episodic, cheap and easy to simulate, easy to score
ta-reinforcement learning

ocumentation, and tutorials

Deep Reinforcement Learning

Deep Reinforcement Learning Nanodegree Links
Course Description
Cartoons
Intuitive RL: Intro to Advantage-Actor-Critic (A2C)
Cheatsheets
Cheat Sheets for AI, Neural Networks, Machine Learning, Deep Learning & Big Data
C++ Python Cheatsheet
Cheat Sheets for AI, Neural Networks, Machine Learning, Deep Learning & Big Data
Cloud
List of Deep Learning Cloud Service Providers
Tenzar
Floydhub
Seedbank
TensorFlow Hub
Google AutoML
Nimblebox
VectorDash
Conferences
NIPS 2017 – videos
ICML 2018 – Stockholm – July 10-15, 2018
Artificial Intelligence Conference – San Francisco – Sep 5-7, 2018 – \$1895+
Deep Learning Summit - Toronto - October 25-26, 2018 - C\$695+
NIPS 2018 - Montréal - December 3-8, 2018 - SOLD OUT
Deep Learning Summit - San Francisco - January 24 - 25, 2019
Community
DRLNG Students
Research Discussions
How to Lead a Discussion of Scientific Journal Articles
Leading a discussion of a scientific paper
Competitions
Pommerman
Halite
NIPS 2018 Competition Track
OpenAI Retro Contest
Courses
TensorFlow for Deep Learning
Stanford
Thomas Simonini : Deep Reinforcement Learning Course
- An Introduction to Reinforcement Learning
- Diving deeper into Reinforcement Learning with Q-Learning
- An introduction to Q-Learning: reinforcement learning
- Curiosity-Driven Learning made easy Part I
Fellowship AI
UC Berkeley - Deep Reinforcement Learning
- Yousof of DRLND: Imitating Learning
Open Machine Learning Course
Machine Learning and Reinforcement Learning in Finance Specialization
Intro to Deep Learning with PyTorch
Łukasz Kaiser "Deep Learning: The Good, the Bad and the Ugly"
Github
30 Amazing Machine Learning Projects for the Past Year (v.2018)
alpha-zero-general: a clean implementation based on AlphaZero for any game in any framework
An AI agent learning to walk in gym's BipedalWalker environment
Deep Learning with PyTorch
Deep reinforcement learning GPU libraries for NVIDIA Jetson with PyTorch, OpenAI Gym, and Gazebo robotics simulator
Deep reinforcement learning GPU libraries for NVIDIA Jetson with PyTorch, OpenAI Gym, and Gazebo robotics simulator.
Deep RL Arm Manipulation
Distributed evolution
Evostra: Evolution Strategy for Python
GAN Lab
Gibson Environment: Real-World Perception for Embodied Agents
Google: Dopamine is a research framework for fast prototyping of reinforcement learning algorithms.
Highly modularized implementation of popular deep RL algorithms in PyTorch
Large-Scale Study of Curiosity-Driven Learning
Mantra: A high-level, rapid development framework for machine learning projects
PlaidML is the easiest, fastest way to learn and deploy deep learning on any device
Random search, hill climbing, policy gradient for CartPole

' http://bit.ly/drIndlinks '	Please Insert > Row and add new links freely
https://www.udacity.com/course/deep-reinforcement-learning-nanodegree--nd893	
https://hackernoon.com/intuitive-rl-intro-to-advantage-actor-critic-a2c-4ff545978752	
https://becominghuman.ai/cheat-sheets-for-ai-neural-networks-machine-learning-deep-learning-big-data-678c51b4b463	
https://d17h27t6h515a5.cloudfront.net/topher/2018/January/5a4d862b_c-python-cheatsheet/c-python-cheatsheet.pdf	
https://becominghuman.ai/cheat-sheets-for-ai-neural-networks-machine-learning-deep-learning-big-data-678c51b4b463	
https://towardsdatascience.com/list-of-deep-learning-cloud-service-providers-579f2c769ed6	30 providers listed
https://www.tenzar.com/	
https://www.floydhub.com/	Creating a 'Run on FloydHub' Button:
http://tools.google.com/seedbank/	
https://www.tensorflow.org/hub/	
https://cloud.google.com/automl/	Train high-quality custom machine learning models with
https://nimblebox.ai/	
https://vectordash.com/	
https://www.facebook.com/pg/nipsfoundation/videos/?ref=page_internal	
https://medium.com/@jianzhang_23841/a-comprehensive-summary-and-categorization-on-reinforcement-learning-papers-at-ic	
https://conferences.oreilly.com/artificial-intelligence/ai-ca/public/register	https://www.youtube.com/watch?v=0f5ytRgDM7g&list=P
https://www.re-work.co/events/deep-learning-summit-toronto-canada-2018	
https://nips.cc/	https://syncdreview.com/2018/09/04/nips-tickets-sell-o
https://www.re-work.co/events/deep-learning-summit-san-francisco-2019	
https://docs.google.com/spreadsheets/d/1veKepeuZ4pFSLO8CldZ7JSkiym41f-e02HOMCfUIM-c/edit#gid=0	Includes github and blog links
https://docs.google.com/spreadsheets/d/1UneDHaS-G3JJQRYvM_eTGc4PWpJOvBwvrkm-CIBrQ/edit#gid=0	https://youtu.be/hitWL7heZio
https://uvic470ecology.weebly.com/uploads/1/2/4/4/12445281/470_howtoleaddiscussion.pdf	
https://arthropodecology.com/2015/01/21/leading-a-discussion-of-a-scientific-paper/	
https://www.pommerman.com/	
https://halite.io/	
https://nips.cc/Conferences/2018/CompetitionTrack	
https://contest.openai.com/2018-1/	
https://www.udacity.com/course/intro-to-tensorflow-for-deep-learning--ud187	Free class from Google and Udacity
https://online.stanford.edu/courses	
https://simoniniithomas.github.io/Deep_reinforcement_learning_Course/	Udacity alum!
https://medium.freecodecamp.org/an-introduction-to-reinforcement-learning-4339519de419	
https://medium.freecodecamp.org/diving-deeper-into-reinforcement-learning-with-q-learning-c18d0db58efe	
https://medium.freecodecamp.org/an-introduction-to-q-learning-reinforcement-learning-14ac0b4493cc	
https://towardsdatascience.com/curiosity-driven-learning-made-easy-part-i-d3e5a2263359	
https://fellowship.ai/	
http://rail.eecs.berkeley.edu/deeprcourse/	
https://youtu.be/tV1vFtoJCxw	https://drInd.slack.com/archives/C9KU4GN6S/p1535092
https://github.com/Yorko/mlcourse_open	Has not covered RL ... yet
https://www.coursera.org/specializations/machine-learning-reinforcement-finance	
https://www.udacity.com/course/deep-learning-pytorch--ud188	
https://www.youtube.com/playlist?list=PLVjVSqmQgPG-vy8vUHQXnQ7-qhsbKjd9s	
https://medium.mybridge.co/30-amazing-machine-learning-projects-for-the-past-year-v-2018-b853b8621ac7	https://github.com/PAIR-code/facets
https://github.com/suragnair/alpha-zero-general	
https://github.com/alirezamika/bipedal-es	
https://github.com/udacity/DL_PyTorch	
https://github.com/dusty-nv/jetson-reinforcement	
https://github.com/dusty-nv/jetson-reinforcement	
https://github.com/udacity/RoboND-DeepRL-Project	
https://github.com/openai/evolution-strategies-starter	
https://github.com/alirezamika/evotra	
https://github.com/poloclub/ganlab	https://poloclub.github.io/ganlab/
https://github.com/StanfordVL/GibsonEnv	http://gibsonenv.stanford.edu/
https://github.com/google/dopamine	https://ai.googleblog.com/2018/08/introducing-new-fram
https://github.com/ShangtongZhang/DeepRL.git	
https://github.com/openai/large-scale-curiosity	
https://github.com/RJT1990/mantra	https://www.youtube.com/watch?v=w5mwhKCHKX8
https://github.com/plaidml/plaidml	https://techcrunch.com/2018/08/16/intel-buys-deep-lea
https://github.com/kvfrans/openai-cartpole	http://kvfrans.com/simple-algorithms-for-solving-cartpole/

PUBLIC SHARED DOCUMENT - Please respect the audience and don't rename the original (make a copy first!)

https://docs.floydhub.com/guides/run_on_floydhub_button/


minimum effort and machine learning expertise.

['L055Epbe6d5ZASsS-3ZhEWZJAgelwe-lf](#)

[it-in-less-than-12-minutes/](#)

https://storage.googleapis.com/gibson_material/Gibson_CVPR2018.pdf
ework-for-flexible.html

[ring-startup-vertex-ai-to-join-its-movidius-unit](#)



[illegible]

Deep Reinforcement Learning Nanodegree Links
Course Description
RL-Adventure-2: Policy Gradients
rlab is a framework for developing and evaluating reinforcement learning algorithms, fully compatible with OpenAI Gym.
Simple example of DQN for Unity using Keras
Stochastic Ensemble Value Expansion
TensorFlow Reinforcement Learning
TensorForce: A TensorFlow library for applied reinforcement learning
Udacity Deep Reinforcement Learning Nanodegree Playground
Unity Machine Learning Agents Toolkit
Unity Marathon Environments
World models
OpenAI Baselines is a set of high-quality implementations of reinforcement learning algorithms.
- PRE Prioritized Replay
- A2C Asynchronous Actor-Critic
- DDPG Deep Deterministic Policy Gradients
Glossary
NVIDIA Deep Learning Glossary
Google: Machine Learning Glossary
An A.I. Glossary (NY Times)
Hardware
Which GPU(s) to Get for Deep Learning
Infographics
A Complete Guide on Getting Started with Deep Learning in Python
Reinforcement Learning Coach by Intel® AI Lab - Supported Algorithms
Jobs
OpenAI – Machine Learning Engineer
Udacity Alumni Slack
Notebooks
Dopamine Colab (Siraj Raval)
Papers - a few "meta" Paper sites (shaded rows), then alphabetical
Deep Reinforcement Learning
Deep Reinforcement Learning: An Overview
OpenAI: Key Papers in Deep RL
reddit MachineLearning
Arxiv Sanity Preserver
Distill
How to Read and Understand a Scientific Paper: A Step-by-Step Guide for Non-Scientists
A Brief Survey of Deep Reinforcement Learning
A Distributional Perspective on Reinforcement Learning
Asynchronous Methods for Deep Reinforcement Learning
Augmented Random Search (ARS)
Autonomous Driving in Reality with Reinforcement Learning and Image Translation
Benchmarking Deep Reinforcement Learning for Continuous Control
Black-Box Data-efficient Policy Search for Robotics
Combined Reinforcement Learning via Abstract Representations
Continuous Control With Deep Reinforcement Learning
Deep Recurrent Q-Learning for Partially Observable MDPs
Deep Reinforcement Learning
Deep Reinforcement Learning in Portfolio Management
Deep Reinforcement Learning that Matters
Deep Reinforcement Learning with Double Q-learning
Distributed Distributional Deterministic Policy Gradients
Dueling Network Architectures for Deep Reinforcement Learning
Emergence of Locomotion Behaviours in Rich Environments
Evolving simple programs for playing Atari games
Hierarchical Reinforcement Learning with the MAXQ Value Function Decomposition
High-dimensional Continuous Control Using Generalized Advantage Estimation
Human-level control through deep reinforcement learning
Is multiagent deep reinforcement learning the answer or the question? A brief survey
Issues in Using Function Approximation for Reinforcement Learning
Lessons Learned Reproducing a Deep Reinforcement Learning Paper
The Lottery Ticket Hypothesis: Finding Sparse, Trainable Neural Networks
Market Making via Reinforcement Learning
Meta-Gradient Reinforcement Learning
Model-Based Reinforcement Learning via Meta-Policy Optimization

' http://bit.ly/drindlinks '	Please Insert > Row and add new links freely
https://www.udacity.com/course/deep-reinforcement-learning-nanodegree--nd893	
https://github.com/higgsfield/RL-Adventure-2	
https://github.com/rli/rliab	
https://github.com/xkiwilabs/DQN_Unity_Keras	
https://github.com/tensorflow/models/tree/master/research/steve	https://arxiv.org/pdf/1807.01675.pdf
https://github.com/deepmind/trrl/	
https://github.com/reinforceio/tensorforce	
https://github.com/RitwikSaikia/udacity_drln_playground	
https://github.com/Unity-Technologies/ml-agents	
https://github.com/Unity-Technologies/marathon-envs	
https://worldmodels.github.io/	
https://github.com/openai/baselines	Updated A2C and PPO2
https://github.com/openai/baselines/blob/master/baselines/deepq/replay_buffer.py	
https://github.com/openai/baselines/tree/master/baselines/a2c	
https://github.com/openai/baselines/tree/master/baselines/ddpg	
https://www.nvidia.com/content/g/pdfs/nvidia-deeplearning-glossary-llkcm.pdf	
https://developers.google.com/machine-learning/glossary/	
https://www.nytimes.com/2018/10/18/business/an-ai-glossary.html	"Reinforcement Learning: A process where machines lea
http://timdettmers.com/2018/08/21/which-gpu-for-deep-learning/	
https://s3-ap-south-1.amazonaws.com/av-blog-media/wp-content/uploads/2018/08/22.jpg	
https://github.com/NervanaSystems/coach/blob/master/img/algorithms.png	
https://www.linkedin.com/jobs/view/816826907/	
https://colab.research.google.com/drive/1sPOu0mmWh2mqZwNAjNf7qei094O3f1-r#scrollTo=mLcmGG387N5t	
https://arxiv.org/abs/1810.06339	150 pages, major improvements to the overview below, 70 pages including Books Surveys and Reports Cours
https://arxiv.org/pdf/1701.07274.pdf	
https://spinningup.openai.com/en/latest/spinningup/keypapers.html	98 linked papers in 12 categories
https://www.reddit.com/r/MachineLearning/	https://youtu.be/SHTOI0KiZnU
http://www.arxiv-sanity.com/	Serving last 50 000+ papers from cs.[CV CL LG AI NE]/s
https://distill.pub/	Machine Learning Research Should Be Clear, Dynamic
https://www.huffingtonpost.com/jennifer-raff/how-to-read-and-understand-a-scientific-paper_b_5501628.html	https://violentmetaphors.com/2013/08/25/how-to-read-ar
https://arxiv.org/pdf/1708.05866.pdf	
https://arxiv.org/pdf/1707.06887.pdf	
https://arxiv.org/pdf/1602.01783.pdf	A2C A3C
https://arxiv.org/pdf/1803.07055.pdf	https://drind.slack.com/archives/C9KU4GN6S/p1533314
https://arxiv.org/pdf/1801.05299.pdf	
https://arxiv.org/pdf/1604.06778.pdf	
https://arxiv.org/pdf/1703.07261.pdf	
https://arxiv.org/abs/1809.04506	A new paper from Yoshua Bengio
https://arxiv.org/pdf/1509.02971.pdf	
https://arxiv.org/pdf/1507.06527.pdf	
https://arxiv.org/pdf/1810.06339.pdf	Yuxi Li (150 pages)
https://arxiv.org/pdf/1808.09940.pdf	Using PPO and DDPG for portfolio optimization
https://arxiv.org/pdf/1709.06560.pdf	
https://arxiv.org/pdf/1509.06461.pdf	
https://openreview.net/pdf?id=SyZlpzbCb	D4PG starts from the DDPG algorithm and includes a nu
https://arxiv.org/pdf/1511.06581.pdf	
https://arxiv.org/pdf/1707.02286.pdf	
https://arxiv.org/abs/1806.05695	Cartesian Genetic Programming (CGP)
https://arxiv.org/pdf/cs/9905014.pdf	
https://arxiv.org/pdf/1506.02438.pdf	
https://storage.googleapis.com/deepmind-media/dqn/DQNNaturePaper.pdf	https://www.cs.swarthmore.edu/~meeden/cs63/s15/natu
https://arxiv.org/pdf/1810.05587.pdf	
http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.73.3097&rep=rep1&type=pdf	
http://amid.fish/reproducing-deep-rl	
https://arxiv.org/pdf/1803.03635.pdf	KEY PAPER. Explains why initialization matters so much
https://arxiv.org/pdf/1804.04216.pdf	
https://arxiv.org/pdf/1805.09801.pdf	
https://arxiv.org/pdf/1809.05214.pdf	

PUBLIC SHARED DOCUMENT - Please respect the audience and don't rename the original (make a copy first!)

arn to do a new task like humans do — through a system of rewards and punishments"

this blog complement it: <https://medium.com/@yuxili>

[ses](#) | [Tutorials](#) | [Conferences, Journals and Workshops](#) | [Blogs](#) | [Testbeds](#) | [Algorithm Implementations](#)

stat.ML

and Vivid. Distill Is Here to Help.

[nd-understand-a-scientific-paper-2/](#)

622000126

CRAR

<https://youtu.be/tJB1qkC1wWM>

This row is INTENTIONALLY a duplicate. Please do not remove.

<https://github.com/qq303067814/Reinforcement-learning-in-portfolio-management->

https://www.youtube.com/playlist?list=PLFU7BilwAjPDqslL9OLm1z7_RXZA1Jyfi

[ire15a.pdf](#)

h in network performance: lucky subnets.

[illegible]

[illegible]

Deep Reinforcement Learning Nanodegree Links
Course Description
Paper Collection of Multi-Agent Reinforcement Learning (MARL)
Multi-Agent Reinforcement Learning: A Report on Challenges and Approaches
Multi-task Deep Reinforcement Learning with PopArt
Neural Fitted Q Iteration - First Experiences with a Data Efficient Neural Reinforcement Learning Method
Noisy Networks for Exploration
Playing Atari with Deep Reinforcement Learning
Prioritized Experience Replay
Proximal Policy Optimization Algorithms
Quantum Reinforcement Learning
Rainbow: Combining Improvements in Deep Reinforcement Learning
Recurrent Experience Replay In Distributed Reinforcement Learning
Reinforcement Learning Using Quantum Boltzmann Machines
Reinforcement Learning with Long Short-Term Memory (LSTM)
Resilient Computing with Reinforcement Learning on a Dynamical System: Case Study in Sorting
RUDDER: Return Decomposition for Delayed Rewards
Sample Efficient Actor-critic with Experience Replay
Sample-Efficient Reinforcement Learning with Stochastic Ensemble Value Expansion
Simple statistical gradient-following algorithms for connectionist reinforcement learning
Some Papers Don't Reproduce. Should We Care?
State of the Art Control of Atari Games Using Shallow Reinforcement Learning
The Effects of Memory Replay in Reinforcement Learning
Thesis: Insights in Reinforcement Learning: Formal analysis and empirical evaluation of temporal-difference learning algorithm
Tight Performance Bounds on Greedy Policies Based on Imperfect Value Functions
Transparency and Explanation in Deep Reinforcement Learning Neural Networks
Troubling Trends in Machine Learning Scholarship
Tuning Recurrent Neural Networks With Reinforcement Learning
Unity: A General Platform for Intelligent Agents
Variational Discriminator Bottleneck
Patents
Dueling Deep Neural Networks
Training Neural Networks Using a Prioritized Experience Memory
Slides
(10) Principles of Deep RL
(Advances in) Quantum Reinforcement Learning
Algorithmic Machine Learning
Tools
PyTorch
Roboschool
ROS
Unity
.
Videos
Reinforcement learning course - David Silver
Deep Reinforcement learning course - Sergey Levine (UC Berkeley, CS294-112)
Two Minute papers
Deep RL Bootcamp 2017 - Pieter Abbeel
Deep RL Bootcamp Lecture 6: Nuts and Bolts of Deep RL Experimentation
Policy Gradient methods and Proximal Policy Optimization (PPO): diving into Deep RL!
Writingat Reward Functions - Bonsai Gre
Deep Reinforcement Learning in the Enterprise: Bridging the Gap from Games to Industry
Advanced Deep Learning and Reinforcement Learning course at UCL by DeepMind
How to cite in Slack (back ticks prevent preview):

' http://bit.ly/drIndlinks '	Please Insert > Row and add new links freely
https://www.udacity.com/course/deep-reinforcement-learning-nanodegree--nd893	
https://github.com/LantaoYu/MARL-Papers	
https://arxiv.org/abs/1807.09427v1	
https://arxiv.org/pdf/1809.04474.pdf	https://deepmind.com/blog/preserving-outputs-precisely-
http://ml.informatik.uni-freiburg.de/former/_media/publications/riecml05.pdf	
https://arxiv.org/pdf/1706.10295.pdf	
https://arxiv.org/pdf/1312.5602.pdf	DeepMind Technologies (2013)
https://arxiv.org/pdf/1511.05952.pdf	
https://arxiv.org/pdf/1707.06347.pdf	
https://arxiv.org/pdf/0810.3828.pdf	See ' http://bit.ly/drIndtalk ' on August 13
https://arxiv.org/pdf/1710.02298.pdf	Mentioned in Yann LeCun's IJCAI '18 keynote:
https://openreview.net/pdf?id=r1yTjAqYX	
https://1qbit.com/wp-content/uploads/2016/12/1QBit-Research-Paper-%E2%80%93-Quantum-Reinforcement-Learning-%E2%80%93-OpenReview.pdf	See ' http://bit.ly/drIndtalk ' on August 13
https://papers.nips.cc/paper/1953-reinforcement-learning-with-long-short-term-memory.pdf	
https://arxiv.org/abs/1809.09261	Implement sorting with RL
https://arxiv.org/pdf/1806.07857.pdf	
https://arxiv.org/pdf/1611.01224.pdf	
https://arxiv.org/pdf/1807.01675.pdf	
http://www.anw.cs.umass.edu/~barto/courses/cs687/williams92simple.pdf	REINFORCE
https://www.dropbox.com/s/wwwyzck6giue1j1/rpan%20-%20The%20Cost%20of%20Reproducibility.pdf?dl=0	Suggestion: shift the trade-off curve entirely, such that t
https://arxiv.org/abs/1512.01563	
https://arxiv.org/pdf/1710.06574.pdf	https://drind.slack.com/archives/CBFS66ANQ/p1535079
https://hadovanhasselt.files.wordpress.com/2015/12/insights_in_reinforcement_learning_hado_van_hasselt.pdf	
http://www.leemon.com/papers/1993wb2.pdf	
https://arxiv.org/abs/1809.06061	
https://arxiv.org/pdf/1807.03341.pdf	
https://arxiv.org/pdf/1611.02796v3.pdf	
https://arxiv.org/pdf/1809.02627.pdf	
https://arxiv.org/pdf/1810.00821.pdf	
https://patentimages.storage.googleapis.com/9c/c8/12/1b4aa1f80fd680/US20180260689A1.pdf	
https://patentimages.storage.googleapis.com/d3/f0/dd/5439a3c2b82867/US20180260707A1.pdf	
http://www.deeplearningindaba.com/uploads/1/0/2/6/102657286/principles_of_deep_rl.pdf	
http://qtml2017.di.univr.it/resources/Slides/Progress-in-Quantum-Reinforcement-Learning.pdf	
https://pytorch.org/	https://pytorch.org/tutorials/
https://blog.openai.com/roboschool/	Alternative to MuJoCo (Not well maintained)
http://www.ros.org/	
https://unity3d.com/machine-learning/	https://github.com/Unity-Technologies/ml-agents/blob/master/README.md
https://www.youtube.com/playlist?list=PL7-jPKtc4r78-wCZcQn5lqyUWhBZ8fOxT	10 videos 131,681 views Last updated on Feb 23, 2016
https://www.youtube.com/playlist?list=PLkFD6_40KJlznC9CDBVTjAF2oyt8_VAe3	24 Lectures, Fall 2017 Session
https://www.youtube.com/channel/UCbfYPyITQ-7l4upoX8nvctg	
https://sites.google.com/view/deep-rl-bootcamp/lectures	15 videos + slides
https://www.youtube.com/watch?v=8EcdaCk9KaQ&t=0s&index=8	
https://youtu.be/5P7I-xPg8u8	
https://youtu.be/0R3PnJEisqk	
https://youtu.be/GOsUHLr4DKE	
https://www.youtube.com/playlist?list=PLqYmG7hTraZDNJre23vqCGIVpfZ_K2RZs	
' http://bit.ly/drIndlinks '	

[while-adaptively-rescaling-targets/](#)

<https://www.cs.toronto.edu/~vmnih/docs/dqn.pdf>

<https://www.facebook.com/ijcaiecai18/videos/2184672041673770/>

he same amount of researcher time naturally produces more robust and reproducible results.

[502000100](#)

[aster/docs/Getting-Started-with-Balance-Ball.md](#)

[illegible]

[illegible]